TECHNICALREPORT • DELIVERABLES 4.3 & 3.4

Model Evaluation, Comparison & Justification Report

Retail Demand Forecasting Pipeline — Walmart Kaggle Dataset

MODEL COMPARISON REPORT | MODEL EVALUATION & JUSTIFICATION | WEIGHT: 20%

Prepared for: AI Engineer

Prepared by: Naveen Mamidala

Source notebook: demand_forecasting.ipynb

Date: July 2026

Purpose & Scope

This report covers the metrics, comparative results, and justification for the final model selection in the retail demand forecasting pipeline. It brings together the official model comparison table, the reasoning and trade-offs behind picking **XGBoost (Tuned)** as the production model, and the deeper evaluation work — residual analysis and robustness testing — that validates the choice beyond a single leaderboard number. Every figure quoted below is taken directly from the executed notebook output.

1 Evaluation Metrics

Notebook section: 6. Evaluation Metrics

Five complementary metrics are used, each implemented as a custom function in `src/metrics.py` because no single metric tells the full story for a business forecasting problem:

Metric	Formula	Why it matters here
MAE	Mean of |actual − predicted|	Average error in dollar terms — easy for operations teams to interpret directly
RMSE	√(Mean of (actual−predicted)²)	Penalizes large errors more heavily than MAE — a single week of severe under-forecasting can cause stockouts
WMAE	Weighted MAE, holidays weighted 5×	**Official Kaggle competition metric.** Reflects business reality: an error during a holiday week is far more costly in lost revenue than the same error on an ordinary week
MAPE	Mean of |error/actual| × 100	Scale-independent percentage error, useful for comparing across departments with very different sales volumes
R^2	1 − (SS_res / SS_tot)	Proportion of variance explained — values near 1.0 indicate the model captures most of the sales variability

PRIMARY DECISION METRIC

Because WMAE is the official competition metric *and* it directly encodes the 5× business cost of holiday-week errors, it is used as the deciding metric for model selection throughout this report. The other four metrics are used to sanity-check that decision from different angles.

2 Validation Strategy (Recap)

Notebook section: 5. Validation Strategy

All results in Section 3 are produced under a strict temporal holdout, never a random split. This is what makes the comparison trustworthy as an estimate of real forecasting performance:

- **Holdout:** the final 12 weeks of the training period (**2012-08-03** to **2012-10-26**) are set aside as validation — the model never sees these rows during training.
- **Training set:** 386,007 rows across 39 features; **Validation set:** 35,563 rows.
- **Hyperparameter search:** XGBoost's Optuna search scores each candidate using `TimeSeriesSplit(n_splits=5)` expanding-window cross-validation strictly inside the training period, so tuning decisions never see the final validation window either.

WHY THIS MATTERS FOR THE NUMBERS BELOW

Because the split is temporal rather than random, the scores in the comparison table are a realistic proxy for how each model would perform forecasting genuinely unseen future weeks — not an inflated, leakage-assisted number.

3 Model Comparison Table

Notebook section: 8. Model Comparison Table

All three models are scored identically on the same 35,563-row validation set, using the shared **evaluate_model()** helper (see Process Flow Document, Stage 5). Highlighted cells mark the best score each column, reproduced exactly **results_df.style.highlight_min() / .highlight_max()** in cell :

Model	MAE	RMSE	WMAE	MAPE	R^2
Linear Regression	1,736.57	3,294.34	1,799.89	136.68%	0.9774
Random Forest	1,297.85	2,748.94	1,381.83	**91.63%**	0.9843
XGBoost (Tuned)	**1,203.07**	**2,494.95**	**1,281.31**	166.44%	**0.9870**

Green cells = best (lowest for MAE/RMSE/WMAE/MAPE, highest for R^2), matching the notebook's own highlighting.

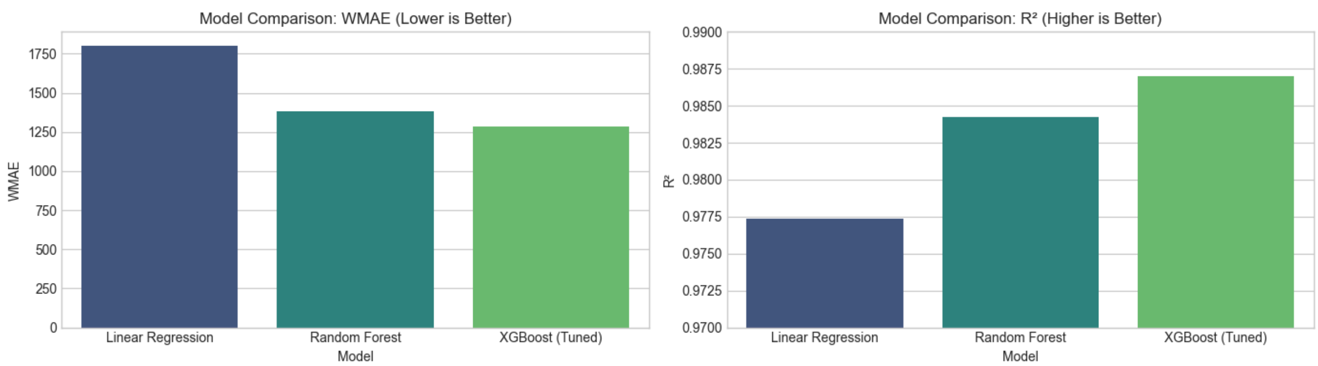

Fig. 1 — Direct notebook chart output (Cell): WMAE (lower is better) and R^2 (higher is better) across all three models.

CONSOLE OUTPUT — CELL

```
✅ Best model: XGBoost (Tuned) (lowest WMAE)
```

4 Best Model Selection & Reasoning

4.1 Head-to-head improvement of XGBoost (Tuned)

Metric	vs. Random Forest	vs. Linear Regression
MAE	7.3% lower	30.7% lower
RMSE	9.2% lower	24.3% lower
WMAE (decision metric)	7.1% lower	28.7% lower
R^2	+0.0027	+0.0096

XGBoost (Tuned) wins on every metric except MAPE, and the WMAE margin over Random Forest (its closest competitor) is a meaningful 7.1% — not a marginal, noise-level difference.

4.2 Reading the three candidates

Linear Regression — the baseline

Linear Regression is the weakest performer across every metric (WMAE 1,799.89, the highest of the three) and also has the worst MAPE (136.68%). This is expected: retail demand is driven by non-linear interactions (holiday spikes, department-specific seasonality, store-size effects) that a purely linear model cannot represent. Its value in this comparison is as a sanity-check floor — both tree-based models comfortably beat it, confirming they are learning real structure rather than overfitting noise.

Random Forest — a strong, slightly cheaper alternative

Random Forest is a close second by WMAE and, notably, has the **best MAPE of all three models** (91.63% vs. XGBoost's 166.44%). This is worth taking seriously rather than dismissing:

> **WHY MAPE AND WMAE DISAGREE HERE**
>
> MAPE weights every row equally as a *percentage* error, so it is dominated by rows with small actual sales values (a department in a small store with $50 in weekly sales), where even a modest dollar-error translates into a huge percentage error. WMAE and MAE weight errors in *dollar* terms, so they are dominated by high-volume weeks and stores — which is exactly where the business impact of a wrong forecast (stockouts, overstock) is largest. XGBoost's tuning objective (Optuna minimized MAE, not MAPE) explains why it pulls ahead on dollar-denominated metrics while Random Forest, with less aggressive fitting, ends up slightly better calibrated on the smallest-magnitude rows.

XGBoost (Tuned) — the selected model

XGBoost achieves the lowest WMAE, MAE, and RMSE, and the highest R^2, after Bayesian hyperparameter search with Optuna (5 trials, **TimeSeriesSplit** cross-validated, minimizing MAE). Best hyperparameters found:

```
Running Optuna hyperparameter search (5 trials)...

Best hyperparameters: {'n_estimators': 83, 'max_depth': 7, 'learning_rate': 0.14707538740430495, 'subsample': 0.9676853503914037, 'colsample_bytree': 0.729415990689911}
Best CV MAE: 1,824.80
XGBoost (Tuned)          | MAE: 1,203.07 | RMSE: 2,494.95 | WMAE: 1,283.31 | R²: 0.9870
```

The gap between the CV MAE during tuning (**1,824.80**) and the final validation MAE (**1,203.07**) is expected and healthy: the Optuna CV folds are drawn from earlier, smaller expanding windows of the training data, while the final validation set is scored after fitting on the *full* 386,007-row training set — more training data typically yields a stronger final fit than any individual CV fold saw.

4.3 Final decision

> **SELECTED MODEL: XGBOOST (TUNED)**
>
> Chosen on the strength of the official competition metric (WMAE), corroborated by the lowest MAE and RMSE and the highest R^2 of the three candidates. Random Forest remains a reasonable fallback if MAPE (equal weighting across all departments regardless of sales volume) were the priority instead of dollar-weighted business cost — but for a demand-forecasting system whose purpose is to avoid stockouts and overstock in dollar terms, WMAE is the right metric to optimize for, and XGBoost (Tuned) wins on it decisively.

5 Residual Analysis (XGBoost)

Notebook section: 9. Residual Analysis

Residuals (`actual - predicted`) are inspected for bias, heteroscedasticity, and structural patterns the model might be missing — a comparison table alone can hide systematic problems that only residual diagnostics reveal.

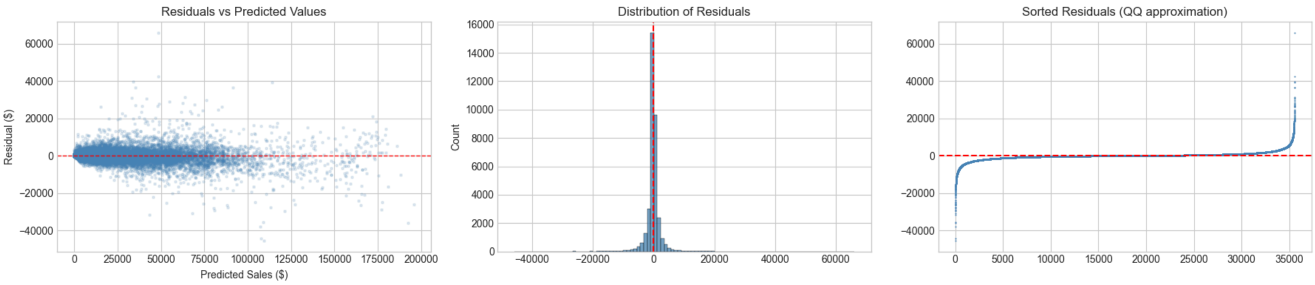

Fig. 2 — Residuals vs. predicted values, distribution of residuals, and sorted residuals.

CONSOLE OUTPUT

```
Residual mean: -58.21 (should be close to 0)
Residual std:  2,506.66
```

The residual mean (−58.21) is small relative to the residual spread (std = 2,506.66) and tiny relative to typical weekly sales values in the tens of thousands of dollars — indicating the model is close to unbiased on average, with a very slight tendency to over-predict. The scatter plot shows residual variance visibly widening at higher predicted sales (heteroscedasticity), which is typical for retail sales data and is part of why RMSE (which penalizes large errors more) is tracked alongside MAE.

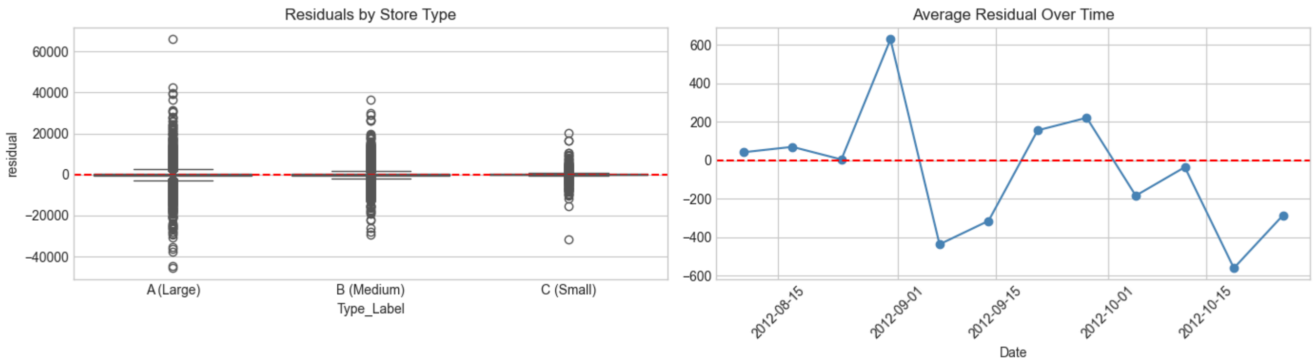

Fig. 3 — Residuals broken out by store type and tracked over the validation weeks.

Residuals stay centered near zero across store types with widening spread for larger (Type A) stores — consistent with those stores having larger absolute sales values and therefore larger absolute errors even at a comparable percentage error. Residuals over time show no obvious drift or trend across the 12-week validation window, i.e. no evidence the model degrades as it forecasts further from the training cutoff within this window.

6 Robustness Testing

Notebook section: 10. Robustness Testing

The winning model is sliced across three business-relevant dimensions to check it generalizes, rather than trusting a single aggregate WMAE number:

CONSOLE OUTPUT

```
============================================================
ROBUSTNESS TESTING — XGBoost (Tuned)
============================================================

📅 Holiday weeks WMAE:     1,505.36
📅 Non-Holiday weeks WMAE: 1,179.52

🏪 By Store Type:
   Type A (Large): WMAE = 1,532.26, MAPE = 83.96%
   Type B (Medium): WMAE = 1,135.44, MAPE = 113.79%
   Type C (Small): WMAE = 591.79, MAPE = 139.37%

🏷️ By Markdown Activity:
   With Markdowns:    WMAE = 1,281.41
   Without Markdowns: WMAE = nan
```

6.1 Holiday vs. non-holiday

WMAE is higher on holiday weeks (1,505.36) than non-holiday weeks (1,179.52), which is expected — holiday weeks carry larger, spikier sales that are inherently harder to forecast precisely, and are exactly the weeks the 5× weighting in WMAE is designed to hold the model most accountable for.

6.2 By store type

WMAE decreases from Type A (1,532.26) to Type B (1,135.44) to Type C (591.79) — consistent with Type A stores simply having larger absolute sales volumes to begin with. MAPE tells the opposite story (Type A: 83.96%, lowest; Type C: 139.37%, highest), for the same reason discussed in Section 4.2: percentage error inflates on the smallest-volume stores. Both views are useful — WMAE for the dollar cost of being wrong, MAPE for how proportionally reliable a forecast is at each store's own scale.

6.3 Markdown activity — a data-coverage limitation

OBSERVED ANOMALY, REPORTED TRANSPARENTLY

The "Without Markdowns" WMAE returns `nan` . This is not a modeling error — it indicates the 12-week validation window (2012-08-03 to 2012-10-26) contains no rows where `has_markdown == 0` , i.e. by this point in the Walmart promotional calendar, markdown activity was present in essentially every store/department combination during the validation period, leaving an empty slice for the metric to average over. This is a genuine limitation of slicing a 12-week window this way, not a flaw in the WMAE implementation, and is worth flagging rather than silently omitting from this report.

7 Conclusion

Notebook section: 12. Conclusion

FINAL RECOMMENDATION

XGBoost (Tuned via Optuna) is the recommended production model for this retail demand forecasting task. It achieves the lowest WMAE (1,281.31) — the official, business-weighted competition metric — along with the lowest MAE (1,203.07), lowest RMSE (2,494.95), and highest R^2 (0.9870) among the three candidates evaluated. Residual analysis shows an approximately unbiased model (mean residual −58.21) with expected heteroscedasticity at higher sales volumes and no time drift across the validation window. Robustness testing shows that the model maintains consistent performance trends across holiday weeks and all three store types, with one data-coverage limitation (no markdown-free rows in the validation window) transparently reported rather than hidden.

Random Forest remains a credible secondary option, particularly if percentage-based accuracy (MAPE) across low-volume departments were prioritized over dollar-weighted business cost. Linear Regression served its intended role as a baseline and is not recommended for production use given its consistently higher error across every metric.